

A multilevel preconditioner for data assimilation with 4D-Var

Alison Ramage*, Kirsty Brown*, Igor Gejadze†

* Department of Mathematics and Statistics, University of Strathclyde

† IRSTEA, Montpellier, France



Data assimilation

- Numerical weather prediction is an IVP: given initial conditions, forecast atmospheric evolution. This needs an accurate initial state.
- Data assimilation uses observations and a short range forecast (the background) to update the model state, and determine a “best possible” atmospheric state.
- The update step, which is chosen to minimise a specific cost function, is known as the analysis.
- Other application areas include hydrology, oceanography, environmental science, data analytics, sensor networks. . .

Physical model

- Evolution equation:

$$\begin{aligned}\frac{\partial \varphi(t)}{\partial t} &= F(\varphi(t)) + f(t), \\ \varphi(0) &= u,\end{aligned}$$

$$u \in X, \quad t \in (0, T), \quad f, \varphi \in Y = L_2(0, T; X)$$

true initial state	\bar{u}
true state evolution	$\bar{\varphi}$
observation operator	$C_{obs} : Y \rightarrow Y_{obs}$
observation error	ξ_o
observations	$\varphi_{obs} = C_{obs}\bar{\varphi} + \xi_o$
background error	ξ_b
background function	$u_b = \bar{u} + \xi_b$

Variational data assimilation problem

- represent model in operator form via control-to-state mapping

$$\varphi = R^{cts}(u)$$

- assume errors ξ_o , ξ_b are normal, unbiased and mutually uncorrelated with positive definite covariance operators

$$V_b(\cdot) = E[\langle \cdot, \xi_b \rangle_X \xi_b], \quad V_o(\cdot) = E[\langle \cdot, \xi_o \rangle_{Y_{obs}} \xi_o]$$

- DA problem: find $v \in X$ which minimises

$$J(v) = \frac{1}{2} \langle V_b^{-1} v, v \rangle_X + \frac{1}{2} \langle V_o^{-1} C_{obs} R^{cts}(u) v, C_{obs} R^{cts}(u) v \rangle_{Y_{obs}}$$

Hessian operator

- define associated **tangent linear operator**

$$R'(u)w = \lim_{\tau \rightarrow 0} \frac{R^{cts}(u + \tau w) - R^{cts}(u)}{\tau}, \quad \forall w \in X$$

and adjoint

$$\langle w, R'^*(u)w^* \rangle_X = \langle R'(u)w, w^* \rangle_Y, \quad \forall w \in X, \forall w^* \in Y$$

- Hessian of DA problem:

$$\mathcal{H}(u) = V_b^{-1} + R'^*(u)C_{obs}^* V_o^{-1} C_{obs} R'(u)$$

- Represent functions using a finite-dimensional basis.
- Rewrite as an **unconstrained** minimisation problem using Lagrange's method.
- Incremental approach: **linearise** evolution operator and solve linearised problem iteratively.
- Require a discrete version of the **tangent linear model** (TLM) and its **adjoint**.
- Each iteration requires one **forward** solution of the TLM equations and one **backward** solution of the adjoint equations.

Hessian matrix

- Hessian of the cost function:

$$\mathcal{H} = V_b^{-1} + R^T C_{obs}^T V_o^{-1} C_{obs} R$$

- Discrete **tangent linear operator** R and its adjoint.
- \mathcal{H} is often too large to be stored in memory.
- Action of **applying \mathcal{H} to a vector** is available, but expensive:
 - involves both **forward** and **backward** solves with the linearised evolution operator and its adjoint.

Approximating the inverse Hessian

Why approximate \mathcal{H}^{-1} ?

- \mathcal{H}^{-1} represents an approximation of the **Posterior Covariance Matrix** (PCM).
- The PCM can be used to find **confidence intervals** and carry out *a posteriori* error analysis.
- $\mathcal{H}^{-1/2}$ can be used in **ensemble forecasting**.
- \mathcal{H}^{-1} , $\mathcal{H}^{-1/2}$ can be used for **preconditioning** in a Gauss-Newton method (focus of this talk).

AIM: construct a **limited-memory approximation** to \mathcal{H}^{-1} using only matrix-vector multiplication.

Return to 4D-Var

- Linear system (within a Gauss-Newton method):

$$\mathcal{H}(\mathbf{u}_k)\delta\mathbf{u}_k = \mathbf{G}(\mathbf{u}_k)$$

Hessian of the cost function $\mathcal{H}(\mathbf{u}_k)$
gradient of the cost function $\mathbf{G}(\mathbf{u}_k)$

- Solve using **P**reconditioned **C**onjugate **G**radient iteration (needs only $\mathcal{H}\mathbf{v}$).
- Convergence depends on eigenvalues of the Hessian

$$\mathcal{H} = V_b^{-1} + R^T C_{obs}^T V_o^{-1} C_{obs} R.$$

- Evaluating $\mathcal{H}\mathbf{v}$ is very expensive, so we need a good preconditioner.

First level preconditioning

- Use the background covariance matrix V_b .
- Projected Hessian:

$$H = (V_b^{1/2})^T \mathcal{H} V_b^{1/2} = I + (V_b^{1/2})^T R^T C_{obs}^T V_o^{-1} C_{obs} R V_b^{1/2}$$


- Easy to recover \mathcal{H} in the original space.
- Eigenvalues of H are usually **clustered** in a narrow band above one, with few eigenvalues distinct enough to contribute noticeably to the Hessian.

[HABEN ET AL., COMPUTERS & FLUIDS 46 (2011)]

- This makes H amenable to **limited-memory approximation**.


Correlation matrix

- \mathcal{H}^{-1} (scaled to have unit diagonal)



Preconditioned correlation matrix

- H^{-1} (after first level preconditioning)



Limited-memory approximation

- Find n_e leading eigenvalues and orthonormal eigenvectors using the **Lanczos** method.
- Construct approximation


$$H \approx I + \sum_{i=1}^{n_e} (\lambda_i - 1) \mathbf{u}_i \mathbf{u}_i^T$$

- Easy to evaluate matrix powers:

$$H^p \approx I + \sum_{i=1}^{n_e} (\lambda_i^p - 1) \mathbf{u}_i \mathbf{u}_i^T$$

Second level preconditioning

- **IDEA**: Construct a **multilevel** approximation to H^{-1} based on a sequence of nested grids.
- Discretise evolution equation on a grid with $m + 1$ nodes (level 0) to represent Hessian H_0 .
- Grid level k contains $m_k = m/2^k + 1$ nodes.



- Identity matrix I_k on grid level k .

Grid transfers with “correction”

- Grid transfer based on piecewise cubic splines:
 - Restriction matrix R_c^f from $k = f$ to $k = c$.
 - Prolongation matrix P_f^c from $k = c$ to $k = f$.
- Construct new operators which transfer a matrix between a coarse grid level c and a fine grid level f .
 - From coarse to fine:

$$M_{c \rightarrow f} = P_f^c(M_c - I_c)R_c^f + I_f$$

- From fine to coarse:

$$M_{f \rightarrow c} = R_c^f(M_f - I_f)P_f^c + I_c$$

Outline of multilevel concept

Given a symmetric positive definite operator H_0 available on the finest grid level in matrix-vector product form:

- 1 restrict H_0 to the **coarsest** grid level (denoted by $H_{0 \rightarrow k}$);
- 2 use a **local preconditioner** to improve the eigenvalue distribution;
- 3 build a **limited-memory approximation** to its inverse square root using the **Lanczos** method (which will form the basis of the local preconditioner at the next level up);
- 4 move up one grid level and repeat.

Multilevel algorithm to approximate H^{-1}

- Represent H_0 at a given level (k , say):

$$H_{0 \rightarrow k} = R_k^0 (H_0 - I_0) P_0^k + I_k$$

- Use **local preconditioner** to improve eigenvalue spectrum:

$$\tilde{H}_{0 \rightarrow k} = (B_k^{k+1})^T H_{0 \rightarrow k} B_k^{k+1}$$

- Find n_k eigenvalues/eigenvectors of $\tilde{H}_{0 \rightarrow k}$ using the Lanczos method.

- Approximate $\tilde{H}_{0 \rightarrow k}^{-1/2}$:

$$\tilde{H}_{0 \rightarrow k}^{-1/2} \approx I_k + \sum_{i=1}^{n_k} \left(\frac{1}{\sqrt{\lambda_i}} - 1 \right) \mathbf{u}_i \mathbf{u}_i^T$$

Local preconditioners

- B_k^{k+1} , constructed on level $k + 1$ and applied on level k .
- On coarsest grid, level $k + 1$ does not exist so set $B_k^{k+1} = I_k$.
- For other levels, construct preconditioners recursively:

$$B_k^{k+1} = \left[B_{k+1}^{k+2} \tilde{H}_{0 \rightarrow k+1}^{-1/2} \right]_{\rightarrow k}, \quad B_k^{k+1 T} = \left[\tilde{H}_{0 \rightarrow k+1}^{-1/2} B_{k+1}^{k+2 T} \right]_{\rightarrow k}$$

- Example: 4 levels ($k = 0, 1, 2, 3$)

$$B_4^5 = I_4,$$

$$B_3^4 = [\tilde{H}_{0 \rightarrow 4}^{-1/2}]_{\rightarrow 3},$$

$$B_2^3 = [[\tilde{H}_{0 \rightarrow 4}^{-1/2}]_{\rightarrow 3} \tilde{H}_{0 \rightarrow 3}^{-1/2}]_{\rightarrow 2},$$

$$B_1^2 = [[[\tilde{H}_{0 \rightarrow 4}^{-1/2}]_{\rightarrow 3} \tilde{H}_{0 \rightarrow 3}^{-1/2}]_{\rightarrow 2} \tilde{H}_{0 \rightarrow 2}^{-1/2}]_{\rightarrow 1},$$

$$B_0^1 = [[[[\tilde{H}_{0 \rightarrow 4}^{-1/2}]_{\rightarrow 3} \tilde{H}_{0 \rightarrow 3}^{-1/2}]_{\rightarrow 2} \tilde{H}_{0 \rightarrow 2}^{-1/2}]_{\rightarrow 1} \tilde{H}_{0 \rightarrow 1}^{-1/2}]_{\rightarrow 0}.$$

Finest level

- We already have H_0 , so precondition to obtain

$$\tilde{H}_0 = B_0^{1T} H_0 B_0^1.$$

- Find n_0 eigenvalues/eigenvectors of \tilde{H}_0 using the Lanczos method.
- Approximate \tilde{H}_0^{-1} :

$$\tilde{H}_0^{-1} \approx I_k + \sum_{i=1}^{n_0} \left(\frac{1}{\lambda_i} - 1 \right) \mathbf{u}_i \mathbf{u}_i^T$$

- Recover projected inverse Hessian using

$$H_0^{-1} = B_0^1 \tilde{H}_0^{-1} B_0^{1T}$$

- The main cost comes from using the Lanczos method, which involves the **Hessian vector product** $H_0 \mathbf{v}$.
- All matrices involved are in **limited-memory** form, and are never actually constructed: instead they are applied using multiplication with eigenvalues and eigenvectors only.
- The remaining problem is how to choose the number of eigenvalues used in the limited-memory approximation on each level

$$N_e = (n_0, n_1, \dots, n_c)$$

to balance **accuracy** and **cost**.

Summary of algorithm

- use $N_e = (n_0, n_1, \dots, n_c)$ eigenvalues at each level

```
[ $\Lambda, \mathcal{U}$ ] = mlevd( $H_0, N_e$ )  
for  $k = k_c, k_c - 1, \dots, 0$   
  compute by the Lanczos method  
  and store in memory  
     $\{\lambda_k^i, U_k^i\}, i = 1, \dots, n_k$  of  $\tilde{H}_{0 \rightarrow k}$   
  using preconditioner  $B_k^{k+1}$   
end
```

- storage:

$$\begin{aligned}\Lambda &= [\lambda_{k_c}^1, \dots, \lambda_{k_c}^{n_{k_c}}, \lambda_{k_c-1}^1, \dots, \lambda_{k_c-1}^{n_{k_c-1}}, \dots, \lambda_0^1, \dots, \lambda_0^{n_0}], \\ \mathcal{U} &= [U_{k_c}^1, \dots, U_{k_c}^{n_{k_c}}, U_{k_c-1}^1, \dots, U_{k_c-1}^{n_{k_c-1}}, \dots, U_0^1, \dots, U_0^{n_0}].\end{aligned}$$

Example


- Test using 1D **Burgers' equation** with initial condition

$$f(x) = 0.1 + 0.35 \left[1 + \sin \left(4\pi x + \frac{3\pi}{2} \right) \right], \quad 0 < x < 1$$

- 1D uniform grid with 7 sensors located at 0.3, 0.4, 0.45, 0.5, 0.55, 0.6, and 0.7 in $[0, 1]$.
- Multilevel preconditioning with **four** grid levels:

k	0	1	2	3
grid points	401	201	101	51

Diagonal of H^{-1}



Assessing approximation accuracy

- **Riemannian** distance:

$$\delta(A, B) = \|\ln(B^{-1}A)\|_F = \left(\sum_{i=1}^n \ln^2 \lambda_i \right)^{1/2}$$

- Compare eigenvalues of H^{-1} and \tilde{H}^{-1} on the finest grid level $k = 0$ using


$$D = \frac{\delta(H^{-1}, \tilde{H}^{-1})}{\delta(H^{-1}, I)}$$

- Vary number of eigenvalues chosen on each grid level

$$N_e = (n_0, n_1, n_2, n_3)$$

Eigenvalues of the inverse Hessian

- Exact (blue circles), approximated (red stars)




$$N_e = (64, 0, 0, 0)$$

$$D = 2.98e - 4$$

Eigenvalues of the inverse Hessian

- Exact (blue circles), approximated (red stars)




$$N_e = (8, 0, 0, 0)$$

$$D = 7.71e - 1$$


Eigenvalues of the inverse Hessian

- Exact (blue circles), approximated (red stars)



Eigenvalues of the inverse Hessian

- Exact (blue circles), approximated (red stars)




$$N_e = (0, 0, 29, 6)$$

$$D = 3.39e - 1$$

Fixed memory ratio

- Fixed memory ratio $R = \sum_{k=0}^{k_c} \frac{n_k}{2^k}$



PCG iteration for one Newton step

- measurement units
 - memory: length of vector on finest grid L
 - cost: cost of HVP on finest grid M

Preconditioner	# CG iterations	storage	cost
none	57	0L	57M
MG(400,0,0,0)	1	400L	402M
MG(4,8,16,32)	4	16L	34M
MG(0,8,16,32)	5	12L	14M
MG(0,0,16,32)	8	8L	10M

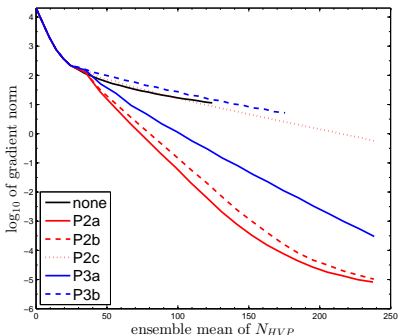

Practical versions

- **Algorithm 1:** apply **mlevd** to H .
 - Still needs a lot of resources for realistic problems.
- **Algorithm 2:** assemble **local** Hessians for each sensor to form H_g , then apply **mlevd** to H_g .
 - Local Hessians could be computed at a coarser grid level l (including discretisation of the tangent linear model). ✓
 - Local Hessians cheaper to compute (use n_l eigenvalues), and can be computed in parallel. ✓
 - Need to store additional local Hessians. ✗
- **Algorithm 3:** apply **mlevd** to local **inverse** Hessians and assemble to form H_g^{rm} , then apply **mlevd** to H_g^{rm} .
 - Needs less storage than Algorithm 2. ✓
 - Can use N_e^l eigenvalues for local Hessians. ✓
 - Introduces an additional layer of approximation. ✗

Iteration counts

Preconditioner	Algorithm	N_e	l	n_l	N_e^l
P2a	2	(200,0,0,0)	1	8	-
P2b	2	(0,8,16,32)	1	8	-
P2c	2	(0,4,8,16)	1	8	-
P3a	3	(0,8,16,32)	1	8	(0,0,8,0)
P3b	3	(0,8,16,32)	2	8	(0,0,0,8)

Number of Hessian vector products



Conclusions and next steps

- Similar results with other configurations (e.g. moving sensors, different initial conditions).
- Multilevel preconditioning looks promising for constructing a good limited-memory approximation to H^{-1} .
- The balance between restrictions on memory/cost limitations may vary between particular applications.
- Identifying globally appropriate values for (n_0, n_1, n_2, n_3) and other parameters is tricky.
- Next steps:
 - extension to higher dimensions;
 - applying the evolution model on coarser grids;
 - application to other sensor networks.

Please join us in June!

27th Biennial Numerical Analysis Conference
University of Strathclyde, Glasgow, Scotland
June 27th-30th 2017



<http://numericalanalysisconference.org.uk/>

Hessian decomposition

- partition domain into subregions and compute **local Hessians** H^l such that

$$H(u) = I + \sum_{l=1}^L (H^l(u) - I)$$

- fewer eigenvalues required for limited-memory representation of each H^l
- local Hessians can be computed in **parallel**
- H^l need not be computed at finest grid level:

$$H_k(u_k) = I_k + \sum_{l=1}^L (H_k^l(u_k) - I_k)$$


- could run local rather than global model

Practical approach: version 1

- Compute limited-memory approximations to **local sensor-based Hessians** on level l using n_l eigenpairs.
- Assemble these to form H_a , then apply **mlevd** to H_a based on a fixed N_e .
- Local Hessians **cheaper to compute**.
- Additional user-specified parameter(s) l , n_l needed.
- **More memory** required as local Hessians must also be stored.

Numerical results

Preconditioner	N_e	l	n_l
P1	(200,0,0,0)	1	8
P2	(0,8,16,32)	1	8
P3	(0,4,8,16)	1	8




$\log_{10}(\text{error})$ vs number of HVP

- Can reduce memory requirements further by using a **multilevel** approximation of each limited-memory local Hessian on level l using n_l eigenpairs.
- Approximate local Hessians by applying **mlevd** to local **inverse** Hessians based on N_e^l .
- Assemble these to form a reduced-memory assembled Hessian H_a^{rm} .
- Use **mlevd** again on H_a^{rm} based on N_e .

Numerical results

Preconditioner	N_e	l	n_l	N_e^l
P1	(200,0,0,0)	1	8	-
P2	(0,8,16,32)	1	8	-
P3	(0,4,8,16)	1	8	-
P4	(0,8,16,32)	1	8	(0,0,8,0)
P5	(0,8,16,32)	2	8	(0,0,0,8)



log₁₀(error) vs number of HVP